

▼ Homework 2 - Distributed Data Managment

Install findspark, pyspark in case it is not installed - if running on Colab.

You can copy the whole notebook to your Google account and work on it on Colab via:

File -> Save a copy in Drive -> Open the copied notebook

```
# * When using the Docker workspace do not run this step *  
IM_RUNNNING_ON_COLAB = True
```

```
if IM_RUNNNING_ON_COLAB:
```

```
    !pip install --force-reinstall pyspark==3.2  
    !pip install findspark
```

Uplaod the data from Moodle, it's a zip file so simply unzip it

```
!unzip /content/random_data.parquet.zip
```

▼ SparkSession is created outside your function

```
import findspark  
findspark.init()  
from pyspark.sql import SparkSession  
import pyspark  
from time import time  
  
def init_spark(app_name: str):  
    spark = SparkSession.builder.appName(app_name).getOrCreate()  
    sc = spark.sparkContext
```

```
return spark, sc
```

```
spark, sc = init_spark('hw2_kmeans')
```

▼ Load samples points

```
data_df = spark.read.parquet("random_data.parquet")
data_df.show(5)
```

You can load the small sample for quick testing and reproducing results:

```
sample_df = spark.read.option("header", True) \
                        .option('inferSchema', True) \
                        .csv('sample_data_84.csv')
sample_df.show(5)
```

▼ Create initials centroids

```
init_centroids = \
    spark.createDataFrame([[6.693, 7.782, 5.63],
                          [3.744, 4.341, 7.225],
                          [9.01, 7.8, 8.03],
                          [2.134, 1.59, 1.93]])
init_centroids.show()
```

▼ Place your kmeans_fit function here

Don't forget to also add it in a separate .py file named HW2_WET_[ID1]_[ID1]

```
# All of your imports should go here
# You cannot use any premade k-means nor import sklearn...
```

```
def kmeans_fit(data: pyspark.sql.DataFrame,
               init: pyspark.sql.DataFrame,
               ...):
```

```

K: int = 4,
max_iter: int = 10):

# imports
from pyspark.sql import SparkSession
from pyspark.ml.feature import VectorAssembler
import pyspark.sql.functions as F
spark = SparkSession.builder.getOrCreate()

def check_convergence(prev_centroids, centroids):
    for i in range(len(centroids)):
        if list(prev_centroids[i]) != list(centroids[i]):
            return False
    return True

# initialization:
columns = data.columns
converged = False
data_points = data.select("*").persist()
# add a column with number 2 to speed things up late
data_points = data_points.withColumn("pow", F.lit(2))
# create a list of centroids
new_centroids = [list(centroid) for centroid in init.collect()]
# print(centroids)
for iteration in range(max_iter):
    if converged:
        break
    else:
        prev_centroids = new_centroids
        distances = []
        for i, centroid in enumerate(prev_centroids):
            distances.append(f'dist_from_c{i}')
            # for each centroid calculate  $\sigma\{(x_j - c_j)^2\}$  by creating  $(x_j - c_j)^2$  column for each j
            for j in range(len(centroid)):
                data_points = data_points.withColumn(f"x_{j}_minus_c_{j}_pow", F.expr(f"{columns[j]} - {centroid[j]}"))\
                    .withColumn(f"x_{j}_minus_c_{j}_pow", F.pow(f"x_{j}_minus_c_{j}_pow", "pow"))
            # now sum  $(x_j - c_j)^2$  columns to get euclidian distance from centroid i
            # no need for sqrt since its monotonically increasing and distances are strictly positive
            data_points = data_points.withColumn(f'dist_from_c{i}', sum([F.col(f"x_{k}_minus_c_{k}_pow") for k in range(len(centroid))]))

        # create a column with the assigned centroid for each point
        cond = F.expr("CASE " + " ".join([f"WHEN {c} = minimum THEN '{i}'" for i, c in enumerate(distances)]) + " END")

```

```

data_points = data_points.withColumn('minimum', F.least(*distances)).withColumn("centroid_id", cond)

# compute new centroids by averaging the points per centroid
new_centroids = data_points.withColumn("centroid_id",F.col("centroid_id").cast("int"))\
    .groupBy("centroid_id").avg(*columns).orderBy("centroid_id")

#
new_centroids = [list(centroid[1:]) for centroid in new_centroids.collect()]
# check convergence
if iteration > 1:
    converged = check_convergence(prev_centroids, new_centroids)

centroids = spark.createDataFrame(new_centroids)
vecAssembler = VectorAssembler(inputCols=centroids.columns,outputCol="centroids")
return vecAssembler.transform(centroids).select("centroids")

```

▼ Test your function output and run time

```

start_time = time()
out = kmeans_fit(data_df, init_centroids)
end_time = time()

print('Final results:')
out.show(truncate=False)
print(f'Total runtime: {end_time-start_time:.3f} seconds')

```

Final results:

```

+-----+
|centroids|
+-----+
|[6.500257701999981,6.499862152000027,6.500300249000017]|
|[4.500187320000003,4.500350787999989,4.500139439999997]|
|[8.499745406999969,8.50008152299997,8.49965887499999]|
|[1.5000080700000005,1.499990830999997,1.500135027500003]|
+-----+

```

Total runtime: 4.384 seconds

▼ Expected results

For the given initialization centroids and the sample_df_84 you got on Moodle, the expected results is:

```
# +-----+
# |centroid|
# +-----+
# |[4.496670602125162,4.495811688311686,4.50180401416766]|
# |[1.4891561106155216,1.5075798394290807,1.4981257805530763]|
# |[8.501673279603231,8.490239925604461,8.505297582145058]|
# |[6.496611142694714,6.508118249005111,6.510762933484945]|
# +-----+
```

```
%%shell
```

```
jupyter nbconvert --to HTML DDM_HW2_STARTER_318155843_206567067.ipynb
```

```
[NbConvertApp] Converting notebook DDM_HW2_STARTER_318155843_206567067.ipynb to HTML
```

```
[NbConvertApp] Writing 606590 bytes to DDM_HW2_STARTER_318155843_206567067.html
```

Notice that the algorithm is deterministic, and we expect your results to be the same for at least 3 decimal numbers after the point.

The ordering of the centroids in the Dataframe and the title may be different.

Don't forget to run your function on the WHOLE data and show the results in the notebook's PDF and HTML :)

Also don't forget to also add your function in a separate .py file named HW2_WET_[ID1]_[ID1]

✓ 1s completed at 7:28 PM



Could not connect to the reCAPTCHA service. Please check your internet connection and reload to get a reCAPTCHA challenge.